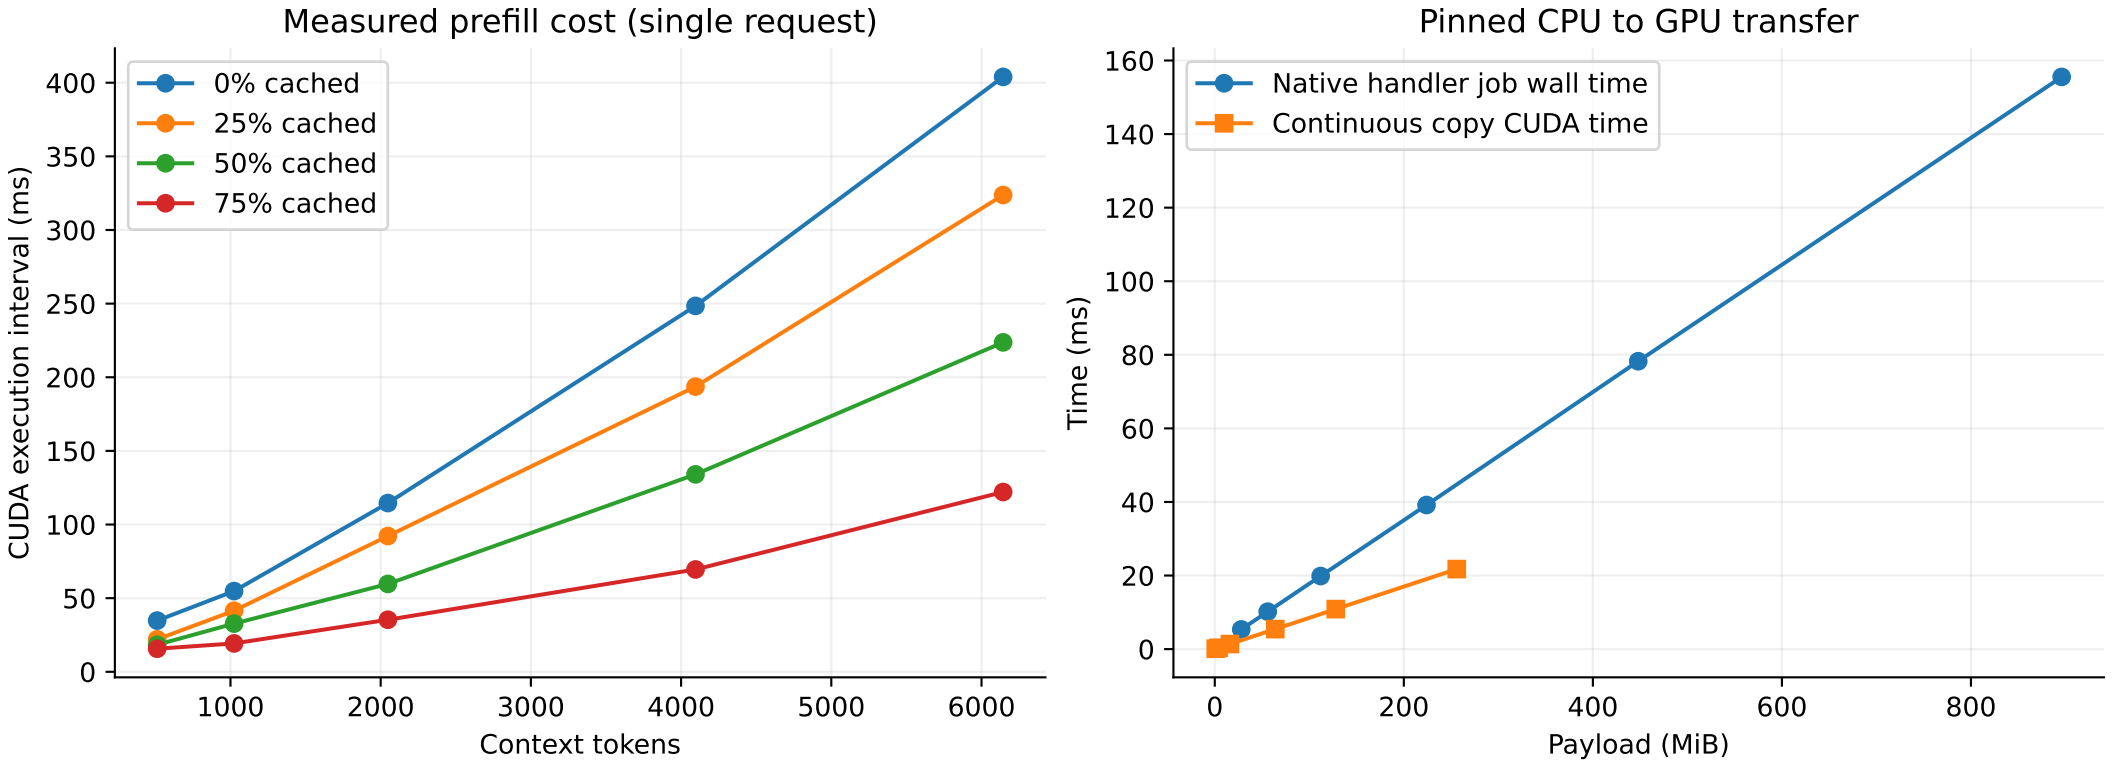
Measured prefill cost (single request)
CUDA execution interval (ms)
Context tokens
0% cached
25% cached
50% cached
75% cached
0
50
100
150
200
250
300
350
400
1000
2000
3000
4000
5000
6000
Pinned CPU to GPU transfer
Time (ms)
Payload (MiB)
Native handler job wall time
Continuous copy CUDA time
0
20
40
60
80
100
120
140
160
0
200
400
600
800