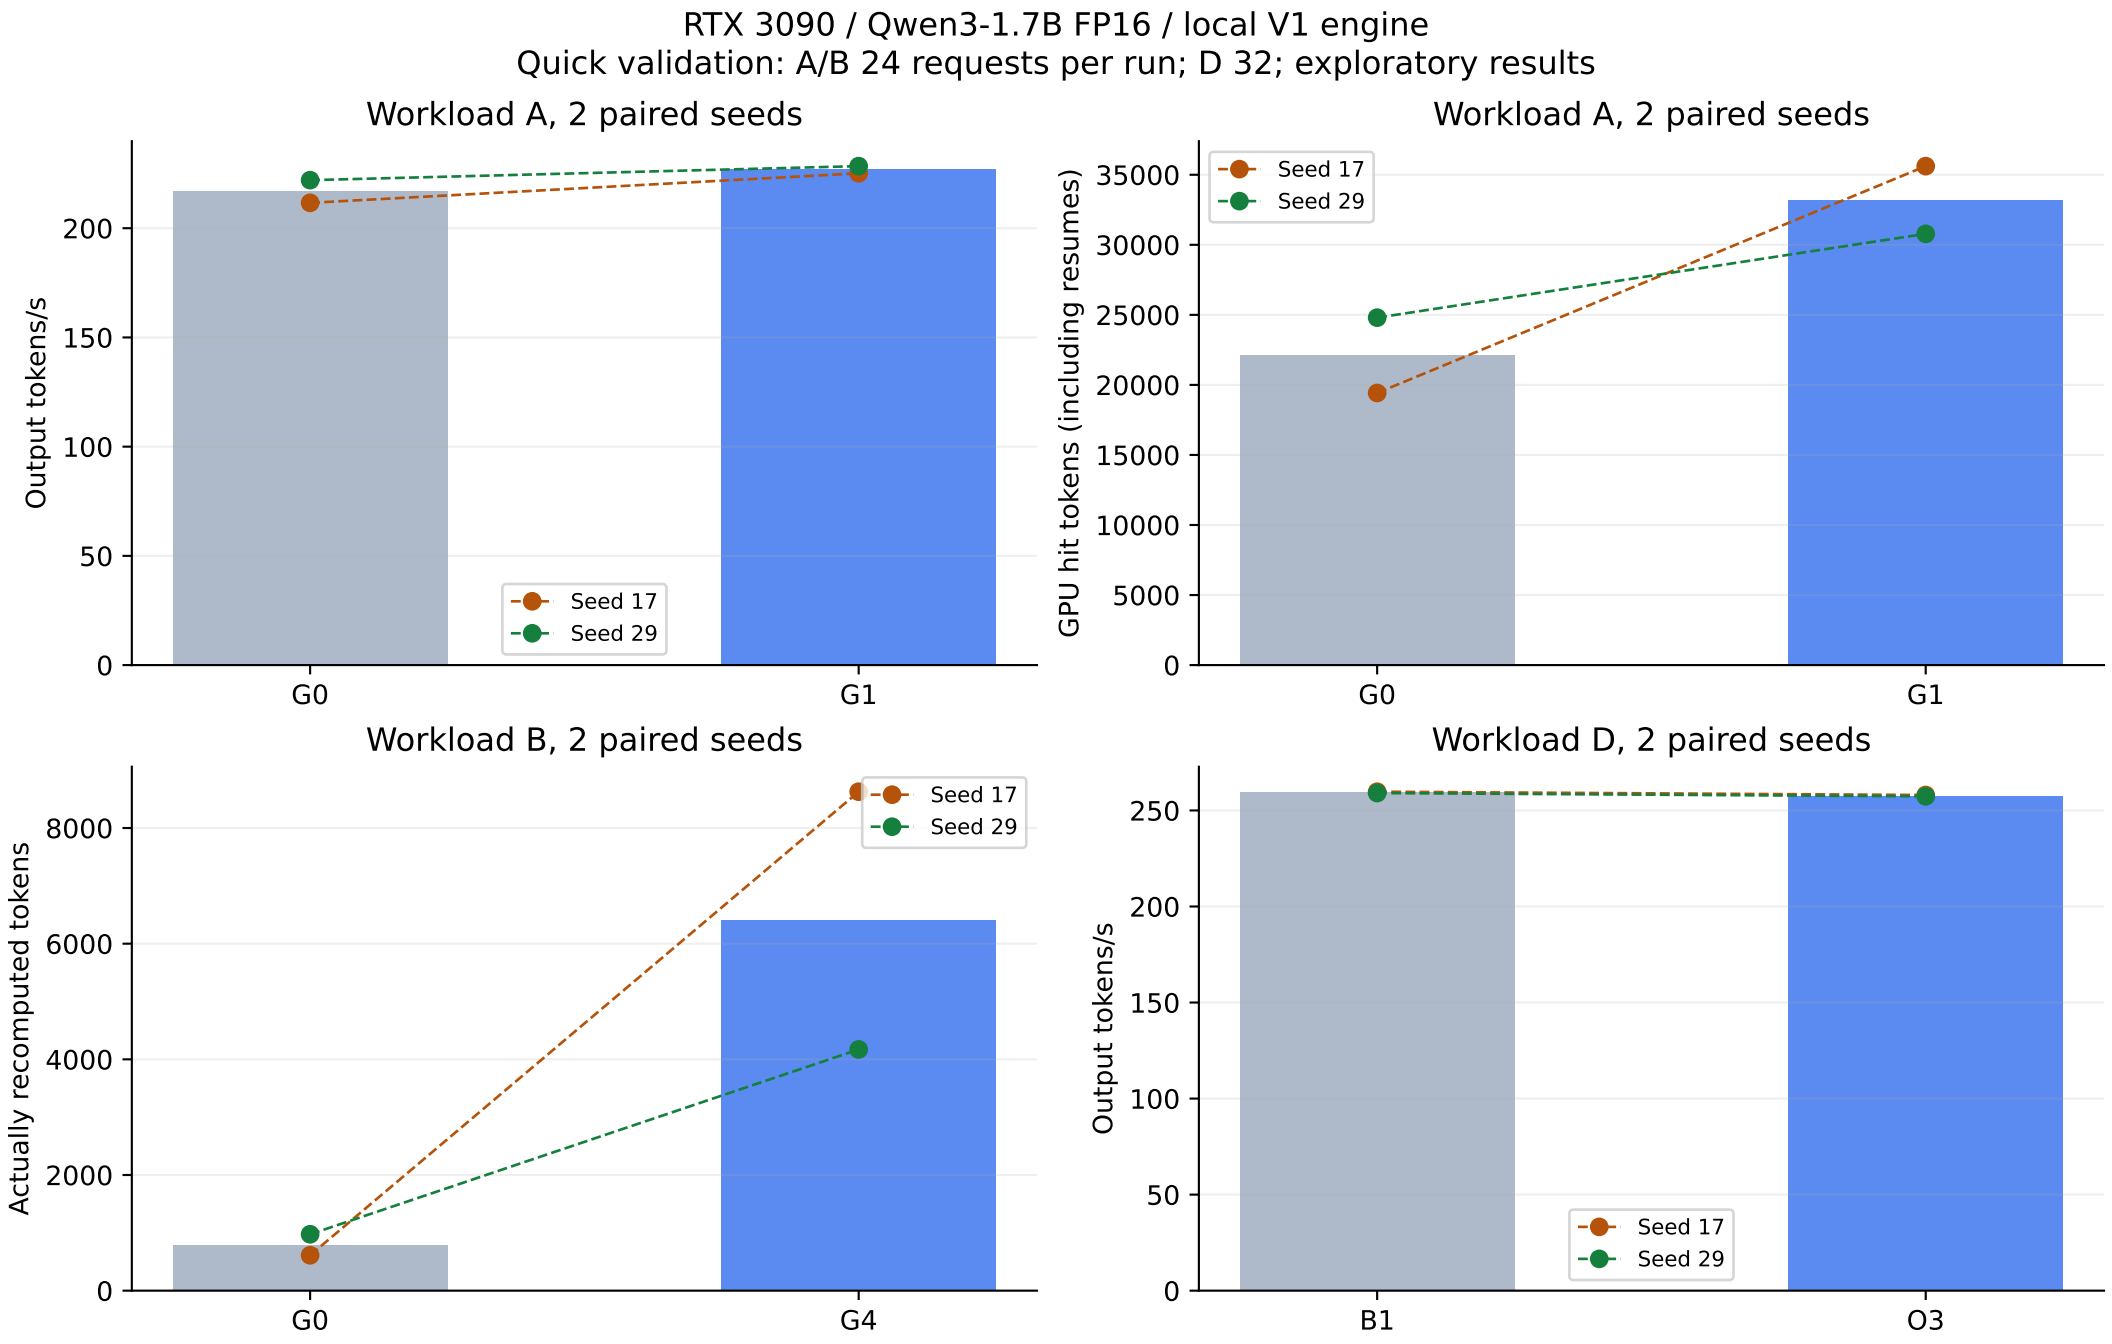
RTX 3090 / Qwen3-1.7B FP16 / local V1 engine
Quick validation: A/B 24 requests per run; D 32; exploratory results
Workload A, 2 paired seeds
Output tokens/s
0
50
100
150
200
Seed 17
Seed 29
G0
G1
Workload A, 2 paired seeds
GPU hit tokens (including resumes)
0
5000
10000
15000
20000
25000
30000
35000
Seed 17
Seed 29
G0
G1
Workload B, 2 paired seeds
Actually recomputed tokens
0
2000
4000
6000
8000
Seed 17
Seed 29
G0
G4
Workload D, 2 paired seeds
Output tokens/s
0
50
100
150
200
250
Seed 17
Seed 29
B1
O3